

- Using the container id, open sh inside the image

```
docker exec -it <container-id> /bin/sh
```

```
(base) → ~ docker exec -it c189882373c9 /bin/sh  
#
```

- Run `cqlsh`

```
(base) → ~ docker exec -it c189882373c9 /bin/sh  
# cqlsh  
Connected to Test Cluster at 127.0.0.1:9042  
[cqlsh 6.0.0 | Cassandra 4.0.7 | CQL spec 3.4.5 | Native protocol v5]  
Use HELP for help.  
cqlsh>
```

- We are now connected to the Cassandra image inside docker.

Using Real-world dataset to get started with the assignment:

1. Download the `cubnb-dataset.txt` from Canvas.
2. Once we have the cassandra image running and are connected to the `sh terminal` inside the image, open the `cubnb-dataset.txt` in a notepad/textpad and copy all the contents (Ctrl + C) and paste the contents in the `sh terminal` (Ctrl + V)
3. Ensure after the completion of all the commands, you get the following output.

```

cqlsh:cubnb> select count(*) from cubnb.property;

count
-----
200

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh:cubnb> select count(*) from cubnb.booking;

count
-----
1500

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh:cubnb> select count(*) from cubnb.tenant;

count
-----
100

(1 rows)

Warnings :
Aggregation query used without partition key

```

Note: In case of an error, please ensure you have spelt the keyspace name and the table names correctly. Keyspace Name: `cubnb`
Table Names: `tenant`, `property`, `booking`

Assignment Submission Instructions:

- Go through below basic operations of Cassandra before starting with Sections 1 and 2.
- Create operation commands and queries for Section 1 & 2. Please attach the screenshots (of query and output) of both the sections.
- Submit the document as PDF file on Canvas.

Note: You may work with a partner on completing this assignment. However, this is an individual assignment; each one must submit your own final deliverable for this assignment.

Basic Operations of Cassandra:

1. **Creating a KeySpace:** A keyspace in Cassandra is a namespace that defines data replication on nodes.

Syntax:

```
CREATE KEYSPACE "KeySpace Name"
```

```
WITH replication = {'class': 'Strategy name', 'replication_factor' : 'Number of replicas'};
```

Example:

```
CREATE KEYSPACE cubnb  
WITH replication = {'class':'SimpleStrategy', 'replication_factor' : 3};
```

2. Using a Keyspace : You can use a created KeySpace using the keyword **USE**

Syntax:
USE <keyspace-name>

Example:
cqlsh> USE cubnb;

3. Creating tables in keyspace :

Syntax:
CREATE (TABLE | COLUMNFAMILY) <tablename>
('<column-definition>', '<column-definition>') (WITH <option> AND
<option>)

Example:
cqlsh:cubnb> CREATE TABLE cubnb.emp(emp_id int PRIMARY
KEY,
emp_name text, emp_city
text, emp_sal varint,
emp_phone varint
);

4. Describe a KeySpace: Verify the keyspaces and its tables.

Syntax:
desc <keyspace-name>

Example:
cqlsh:cubnb> desc cubnb;

5. Drop a KeySpace: Delete the keyspaces and its tables.

Syntax:
drop keyspace <keyspace-name>

Example:
cqlsh:cubnb> drop keyspace cubnb;

Section 1 :

Create commands for below operations as mentioned above and attach screenshots in the assignment.

1. **Create** a keyspace **cubnb** for the application having class as SimpleStrategy and replication_factor as 1.

Keyspace name: `cubnb`.

2. **Use** the keyspace `cubnb` created above.

3. **Create** a table for **tenant** if not exists. Table name: `tenant`.

Columns: `tenant_id` (int, primary key), `mobile` (bigint), `password` (string), `tenant_name` (string).

4. **Create** a table for **property** if not exists. Table name: `property`.

Columns: `property_id` (int, primary key), `property_name` (string), `host_id` (int, primary key),
`location` (string), `type` (string), `price` (int), `number_of_rooms` (int), `max_occupancy` (int),
`number_of_reviews` (int), (set<string> type), `status` (string).
`amenities`

5. **Create** a table for **booking** if not exists. Table name: `booking`.

Columns: `booking_id` (int, primary key), `property_id` (int, primary key), `total_occupancy` (int),
`tenant_id` (int, primary key), `status` (string), `no_of_days` (int), `start_date` (date), `end_date` (date).

Section 2 :

Create queries for below questions and attach screenshots of the query and output in the assignment.

Note: WE HAVE ALREADY IMPORTED DATASET IN THE ABOVE INSTRUCTIONS SECTION

1. Get the booking_id, property_id, tenant_id and number of days a property is booked in the month of Feb '2022 for tenant-id 80855.
2. Get the count of properties whose reviews are more than 200.
3. List all properties from “Brooklyn” location with type as apartment and have status as available.
4. a) Get all the properties which are handled by host_id 1029.
b) Add new amenity ‘Heater’ to the above results (in a single query).
c) Show the updated results from part b query.
5. a) Add your details in tenant table with tenant_id as 3287
b) Change the password of the tenant_id 3287 from the existing to ‘password@123’ and display the change
6. Remove an amenity ‘GardenView’ for host_id 1012 and property_id 110 and display the results.
7. List the properties which have a ‘BeachView’ in its amenities and whose rent price is more than 1000.
8. Get the number of highest occupancy of properties for 8 rooms.